

Hollow Heaps

Thomas Dueholm Hansen¹ Haim Kaplan²
Robert E. Tarjan³ Uri Zwick²

¹ Department of Computer Science,
Aarhus University, Denmark.

² School of Computer Science,
Tel Aviv University, Israel.

³ Department of Computer Science,
Princeton University, USA.

July 7, 2015

Heap operations

- *make-heap()*: Make an empty heap.
- *insert*(e, k, h): Insert item e with key k in the heap h .
- *find-min*(h): Return the item with minimum key in h .
- *delete-min*(h): Remove the item with minimum key from h .
- *decrease-key*(e, k, h): Decrease the key of item e in h to k .
- *delete*(e, h): Remove the item e from h .
- *meld*(h_1, h_2): Return a heap containing all items in h_1 and h_2 .

Let n be the number of items in the heap.

- **Binary heaps:** *insert*, *decrease-key*, and *delete-min* each take $O(\log n)$ time in the worst case.
- **Fibonacci heaps**, Fredman and Tarjan (1987): Optimal **amortized** time bounds.

<i>insert</i> :	$O(1)$
<i>decrease-key</i> :	$O(1)$
<i>delete-min</i> :	$O(\log n)$
<i>meld</i> :	$O(1)$

- Brodal (1996); Brodal, Lagogiannis, and Tarjan (2012): Optimal **worst-case** bounds.
- Numerous other heap variants have been proposed.


- **Dijkstra's algorithm** (1959) for solving single source shortest paths problems with non-negative weights.
- Algorithms for finding undirected and directed **minimum spanning trees** (Gabow, Galil, Spencer, and Tarjan; 1986).
- Etc.

- **Goal:** Find a **simple** and **practical** data structure with optimal worst-case guarantees.
- The design space is rich, and better algorithms may yet be found.
- Fibonacci heaps get enough exposure to motivate further study.


- **Goal:** Find a **simple** and **practical** data structure with optimal worst-case guarantees.
- The design space is rich, and better algorithms may yet be found.
- Fibonacci heaps get enough exposure to motivate further study.
- **Pairing heaps** (Fredman, Sedgwick, Sleator, and Tarjan; 1986) outperform Fibonacci heaps in simplicity and performance.
 - Fredman (1999): *decrease-key* requires $\Omega(\log \log n)$ amortized time.
 - Pettie (2005): $O(2^{2\sqrt{\log \log n}})$ upper bound for *decrease-key*.

Heap-order

- The key of a node is at least as large as the key of its parent.



Linking trees while preserving heap-order



Characteristics of Fibonacci heaps

- A **Fibonacci heap** is a set S of heap-ordered trees.
 - It maintains a pointer to the node with minimum key.

Characteristics of Fibonacci heaps

- A **Fibonacci heap** is a set S of heap-ordered trees.
 - It maintains a pointer to the node with minimum key.
- Each node has a **rank** that matches its degree, and nodes are possibly **marked**.
- *decrease-key* may trigger **cascading cuts**: **Mark** the first unmarked ancestor, and cut all ancestors in between.
 - **4 pointers per node**: Parent, child, left sibling, and right sibling.

Characteristics of Fibonacci heaps

- A **Fibonacci heap** is a set S of heap-ordered trees.
 - It maintains a pointer to the node with minimum key.
- Each node has a **rank** that matches its degree, and nodes are possibly **marked**.
- *decrease-key* may trigger **cascading cuts**: **Mark** the first unmarked ancestor, and cut all ancestors in between.
 - **4 pointers per node**: Parent, child, left sibling, and right sibling.
- The amortized time per *delete-min* is bounded by the maximum rank.
 - By induction, the number of nodes in a tree of rank k is at least $F_{k+1} = 1 + F_1 + F_2 + \dots + F_{k-1} = O(1.618^k)$.
 - The maximum rank is at most $\log_{\varphi} n = O(\log n)$.

We introduce **Hollow heaps**, a simple alternative to **Fibonacci heaps**. The distinguishing characteristics of hollow heaps are:

- Each heap is a single heap-ordered **directed acyclic graph** (DAG) that represents all comparisons explicitly.

We introduce **Hollow heaps**, a simple alternative to **Fibonacci heaps**. The distinguishing characteristics of hollow heaps are:

- Each heap is a single heap-ordered **directed acyclic graph** (DAG) that represents all comparisons explicitly.
- There are **no parent pointers**, and *decrease-key* takes $O(1)$ time worst-case and does no cascading cuts.

We introduce **Hollow heaps**, a simple alternative to **Fibonacci heaps**. The distinguishing characteristics of hollow heaps are:

- Each heap is a single heap-ordered **directed acyclic graph** (DAG) that represents all comparisons explicitly.
- There are **no parent pointers**, and *decrease-key* takes $O(1)$ time worst-case and does no cascading cuts.
- Children are stored in **singly linked lists**.

We introduce **Hollow heaps**, a simple alternative to **Fibonacci heaps**. The distinguishing characteristics of hollow heaps are:

- Each heap is a single heap-ordered **directed acyclic graph** (DAG) that represents all comparisons explicitly.
- There are **no parent pointers**, and *decrease-key* takes $O(1)$ time worst-case and does no cascading cuts.
- Children are stored in **singly linked lists**.
- The heap contains **hollow nodes**, and *delete-min* takes $O(\log N)$ amortized time, where N is the total number of nodes.

We introduce **Hollow heaps**, a simple alternative to **Fibonacci heaps**. The distinguishing characteristics of hollow heaps are:

- Each heap is a single heap-ordered **directed acyclic graph** (DAG) that represents all comparisons explicitly.
- There are **no parent pointers**, and *decrease-key* takes $O(1)$ time worst-case and does no cascading cuts.
- Children are stored in **singly linked lists**.
- The heap contains **hollow nodes**, and *delete-min* takes $O(\log N)$ amortized time, where N is the total number of nodes.
- The structure is **exogenous** rather than endogenous: nodes hold items rather than being items.

Using hollow nodes: deletion in linked lists

Doubly linked list:




Singly linked list:




Using hollow nodes: deletion in linked lists

Doubly linked list:




Singly linked list:




Using hollow nodes: deletion in linked lists

Doubly linked list:



Singly linked list:



- Deletions in a singly linked list create “hollow” nodes.

Using hollow nodes: deletion in linked lists

Doubly linked list:



Singly linked list:



- Deletions in a singly linked list create “hollow” nodes.

Using hollow nodes: deletion in linked lists

Doubly linked list:



Singly linked list:



- Deletions in a singly linked list create “hollow” nodes.
- They can only be removed from nodes pointing to them.

Naive hollow heaps

The heap is referenced by its root.

- $\text{meld}(h_1, h_2)$: Link h_1 and h_2 .
- $\text{insert}(e, k, h)$: Create a node x with key k and item e . Link x and h .
- $\text{find-min}(h)$: Return the item of h .
- $\text{decrease-key}(e, k, h)$: Let x be the node containing e . Move e to a new node y with key k , making x **hollow**. Link y and h .
- $\text{delete-min}(h)$: Delete h and all its immediate hollow descendants. Link the exposed full nodes.

Naive hollow heaps

The heap is referenced by its root.

- $\text{meld}(h_1, h_2)$: Link h_1 and h_2 .
- $\text{insert}(e, k, h)$: Create a node x with key k and item e . Link x and h .
- $\text{find-min}(h)$: Return the item of h .
- $\text{decrease-key}(e, k, h)$: Let x be the node containing e . Move e to a new node y with key k , making x **hollow**. Link y and h .
- $\text{delete-min}(h)$: Delete h and all its immediate hollow descendants. Link the exposed full nodes.

Note: This version is inefficient.

Making hollow heaps efficient

- Assign an integer **rank** to every node.
- Initially nodes have rank 0.
- **Ranked link:** Link two nodes of the same rank and increase the rank of the winner by 1.
- **Unranked link:** Link two nodes without changing their ranks.

Making hollow heaps efficient

- Assign an integer **rank** to every node.
- Initially nodes have rank 0.
- **Ranked link:** Link two nodes of the same rank and increase the rank of the winner by 1.
- **Unranked link:** Link two nodes without changing their ranks.
- We allow hollow nodes to have up to two parents.
- **Heap order:** The key of a node is at least as large as the keys of both its parents.

Hollow heaps


- $\text{meld}(h_1, h_2)$: Link h_1 and h_2 .
- $\text{insert}(e, k, h)$: Create a node x with key k , item e , and rank $x.\text{rank} = 0$. Link x and h .
- $\text{find-min}(h)$: Return the item of h .
- $\text{decrease-key}(e, k, h)$:
 - Let x be the node containing e .
 - Move e to a new node y with key k and rank $y.\text{rank} = \max\{0, x.\text{rank} - 2\}$.
 - Make x **hollow** and make y the **second parent** of x .
 - Link y and h .

Hollow heaps

- $\text{meld}(h_1, h_2)$: Link h_1 and h_2 .
- $\text{insert}(e, k, h)$: Create a node x with key k , item e , and rank $x.\text{rank} = 0$. Link x and h .
- $\text{find-min}(h)$: Return the item of h .
- $\text{decrease-key}(e, k, h)$:
 - Let x be the node containing e .
 - Move e to a new node y with key k and rank $y.\text{rank} = \max\{0, x.\text{rank} - 2\}$.
 - Make x **hollow** and make y the **second parent** of x .
 - Link y and h .
- $\text{delete-min}(h)$:
 - Delete h and all its immediate hollow descendants **that only have one parent**.
 - While possible, link the exposed full nodes with **ranked links**.
 - Link the remaining nodes with **unranked links**.


Example

- The heap after successive insertions of items with keys 14, 11, 5, 9, 0, 8, 10, 3, 6, 12, 13, 4.



Example


- The heap after successive insertions of items with keys 14, 11, 5, 9, 0, 8, 10, 3, 6, 12, 13, 4.



- Next:** *delete-min*


Example

- After *delete-min*.



Example


- After *delete-min*.



- **Next:** Decrease 5 to 1.


Example

- After decreasing 5 to 1.



Example


- After decreasing 5 to 1.



- **Next:** Decrease 3 to 2 and 8 to 7.


Example

- After decreasing 3 to 2 and 8 to 7.



Example


- After decreasing 3 to 2 and 8 to 7.



- Next:** *delete-min*


Example

- After *delete-min*.



Example


- After *delete-min*.



- **Next:** *delete-min*

Example

- After *delete-min*.



- Let N be the number of nodes.
- We prove these **amortized** time bounds for hollow heaps:

insert : $O(1)$

decrease-key : $O(1)$

delete-min : $O(\log N)$

meld : $O(1)$

- N is at most the number of items n plus the number of performed *decrease-key* operations.
- In many applications $N = O(\text{poly}(n))$, and otherwise the bound can be improved by occasionally **rebuilding** the heap.


- **Eager implementation:** No second parents. Nodes are moved immediately by *decrease-key*.
- *decrease-key*(e, k, h):
 - Let x be the node containing e .
 - Move e to a new node y with key k and rank $y.rank = \max\{0, x.rank - 2\}$, making x **hollow**.
 - Move all but the **two highest ranked children** of x to y .
 - Link y and h .
- *delete-min*(h):
 - Delete h and all its immediate hollow descendants.
 - While possible, link the exposed full nodes with **ranked links**.
 - Link the remaining nodes with **unranked links**.

Bounding the maximum rank

- A **full** node of rank k has k ranked children with ranks $0, \dots, k - 1$.
- A **hollow** node of rank k has 2 ranked children of ranks $k - 1$ and $k - 2$ (or fewer if $k \leq 1$).

Bounding the maximum rank

- A **full** node of rank k has k ranked children with ranks $0, \dots, k-1$.
- A **hollow** node of rank k has 2 ranked children of ranks $k-1$ and $k-2$ (or fewer if $k \leq 1$).
- By induction, the number of nodes in a tree of rank k is at least $F_{k+3} - 1 = 1 + (F_{k+2} - 1) + (F_{k+1} - 1)$.




- The maximum rank is at most $\log_{\varphi} N = O(\log N)$.

- **Potential argument:** All **unranked children** and **children of hollow nodes** have 1 unit of potential each.
- The potential pays for **linking** and for **deleting hollow nodes**.
- *insert*: Create 1 new unranked child (+1).
- *decrease-key*: Create 1 new unranked child and 2 new children of a hollow node (+3).
- *delete-min*: Create at most $O(\log N)$ new unranked children.

- **Potential argument:** All **unranked children** and **children of hollow nodes** have 1 unit of potential each.
 - The potential pays for **linking** and for **deleting hollow nodes**.
 - *insert*: Create 1 new unranked child (+1).
 - *decrease-key*: Create 1 new unranked child and 2 new children of a hollow node (+3).
 - *delete-min*: Create at most $O(\log N)$ new unranked children.
-
- The analysis of **lazy hollow heaps** is essentially the same.
 - In this case nodes are assigned **virtual children** corresponding to actual children in the eager version.


Lazy hollow heap, implementation

- **3 pointers per node:** Child, right sibling, and **second parent** (can be replaced by a single bit).
- **2 additional pointers to and from items:** Nodes must hold items rather than be items; the structure is **exogenous** rather than **endogenous**.




Lazy hollow heap, implementation

- **3 pointers per node:** Child, right sibling, and **second parent** (can be replaced by a single bit).
- **2 additional pointers to and from items:** Nodes must hold items rather than be items; the structure is **exogenous** rather than **endogenous**.




Lazy hollow heap, implementation

- **3 pointers per node:** Child, right sibling, and **second parent** (can be replaced by a single bit).
- **2 additional pointers to and from items:** Nodes must hold items rather than be items; the structure is **exogenous** rather than **endogenous**.




Lazy hollow heap, implementation

- **3 pointers per node:** Child, right sibling, and **second parent** (can be replaced by a single bit).
- **2 additional pointers to and from items:** Nodes must hold items rather than be items; the structure is **exogenous** rather than **endogenous**.




Lazy hollow heap, implementation



- **Invariant:** A hollow node is the first child of its second parent.

Lazy hollow heap, implementation



- **Invariant:** A hollow node is the first child of its second parent.
- When traversing a list of children, we reach the end when either:
 - There is no next sibling.
 - The **second parent pointer** points back to the parent whose children we are traversing.
- Second parent pointers also show whether nodes have 1 or 2 parents.

- An **experimental study** of hollow heaps.
- Non-cascading Fibonacci heap without hollow nodes?
- **Pairing heaps** (Fredman, Sedgwick, Sleator, and Tarjan; 1986) outperform Fibonacci heaps in practice, although *decrease-key* requires $\Omega(\log \log n)$ time.
- What can be said about **hollow pairing heaps**?
- Array-based heaps outperform pointer-based heaps in practice, but the worst-case bounds are non-optimal. Can we get the best of both worlds?

- An **experimental study** of hollow heaps.
- Non-cascading Fibonacci heap without hollow nodes?
- **Pairing heaps** (Fredman, Sedgewick, Sleator, and Tarjan; 1986) outperform Fibonacci heaps in practice, although *decrease-key* requires $\Omega(\log \log n)$ time.
- What can be said about **hollow pairing heaps**?
- Array-based heaps outperform pointer-based heaps in practice, but the worst-case bounds are non-optimal. Can we get the best of both worlds?

Thank you for listening!

- Larkin, Sen, and Tarjan (2014) made an **experimental study** of various heaps.
- We used their framework to make some initial tests of hollow heaps with randomly generated benchmarks.
- 2,000,000 insertions and deletions, and 5,000,000 *decrease-key* operations.
 - Fibonacci heap: 2488270 ms
 - Pairing heap: 1796165 ms
 - Lazy hollow heap: 1940643 ms